

Project3

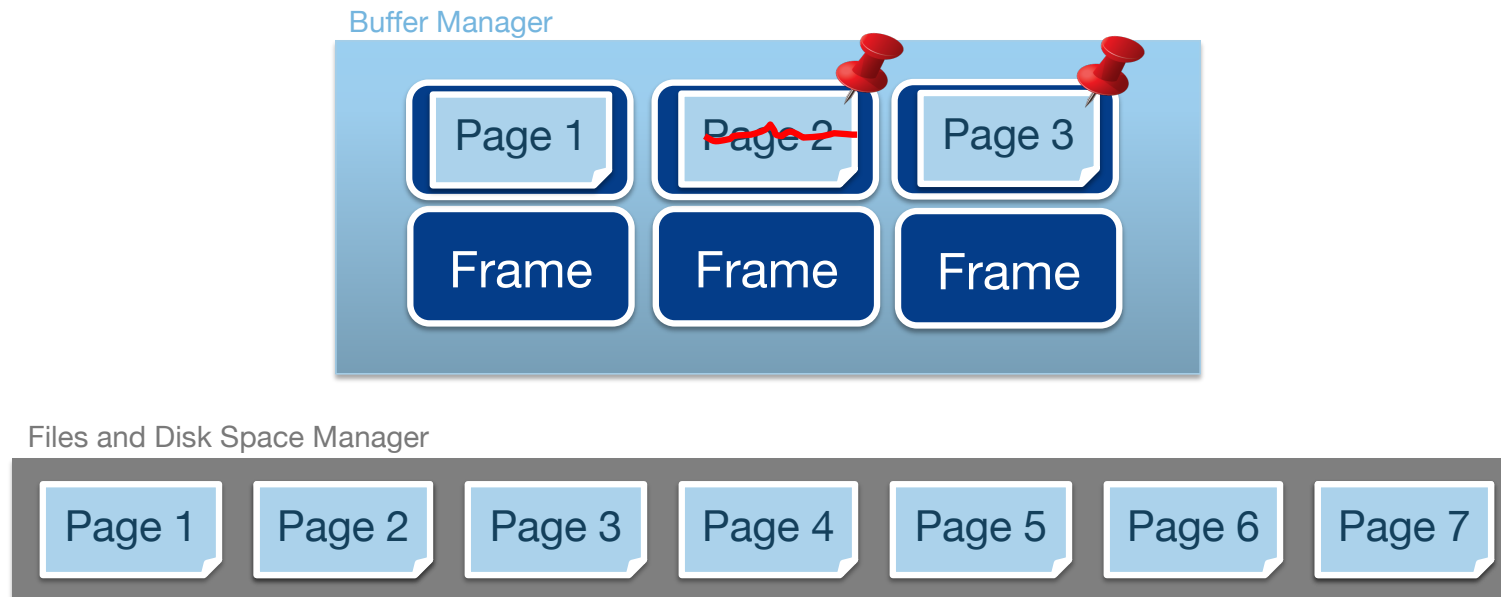
Buffer management

Project Hierarchy

- Your project hierarchy should be like this.
 - Your_hconnect_repo
 - project3
 - include/
 - lib/
 - Makefile
 - src/
- If your Makefile doesn't make libbpt.a library file at the exact path, you'll get zero score. (your_hconnect_repo/project3/lib/libbpt.a)

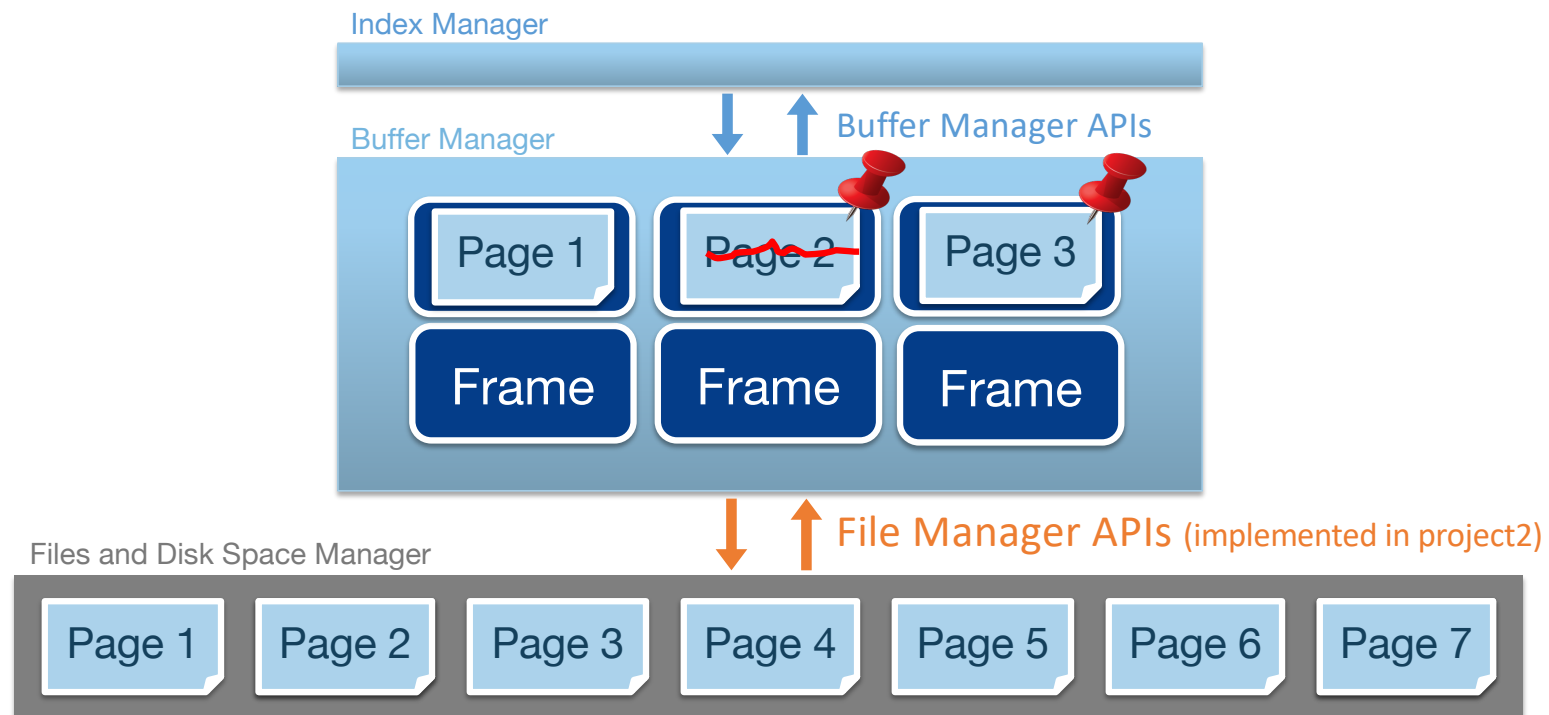
Buffer Management

- Current disk-based b+tree doesn't support buffer management.
- Our goal is to implement **in-memory buffer manager** to caching on-disk pages.



Buffer Management Layer

- File manager APIs should be called only within the buffer manager layer.



Project Specification

- Define the buffer block structure, which must contain at least those fields.
- **Physical frame:** containing up to date contents of target page.
 - **Table id:** the unique id of table (per file)
 - **Page number:** the target page number within a file.
 - **Is dirty:** whether this buffer block is dirty or not.
 - **Is pinned:** whether this buffer block is accessed right now.
 - **LRU list next (prev) :** buffer blocks are managed by LRU list.
 - Other information can be added with your own buffer manager design.

Buffer Structure

frame (page size : 4096 bytes)
table_id
page_num
is_dirty
is_pinned
next/prev of LRU

⋮

Project Specification

➤ Implement database initialization function.

- **int init_db (int num_buf);**
- Allocate the buffer pool (array) with the given number of entries.
- Initialize other fields (state info, LRU info..) with your own design.
- If success, return 0. Otherwise, return non-zero value.

➤ **open_table** interface

- **int open_table (char *pathname);**
- Open existing data file or create one if not existed. You must give the same table id when db opens the same table more than once after init_db(). (the length of pathname ≤ 20)
- If success, return the **unique table id**, which represents the own table in this database. (Return negative value if error occurs)
- You have to maintain a table id once open_table() is called, which is matching file descriptor or file pointer depending on your previous implementation. (table id ≥ 1 and maximum allocated id is set to 10)

Project Specification

- A table id needs to be passed to the index manager APIs to select the table where the operation will be executed.
 - `int db_insert (int table_id, int64_t key, char * value);`
 - `int db_find (int table_id, int64_t key, char* ret_val);`
 - `int db_delete (int table_id, int64_t key);`

Project Specification

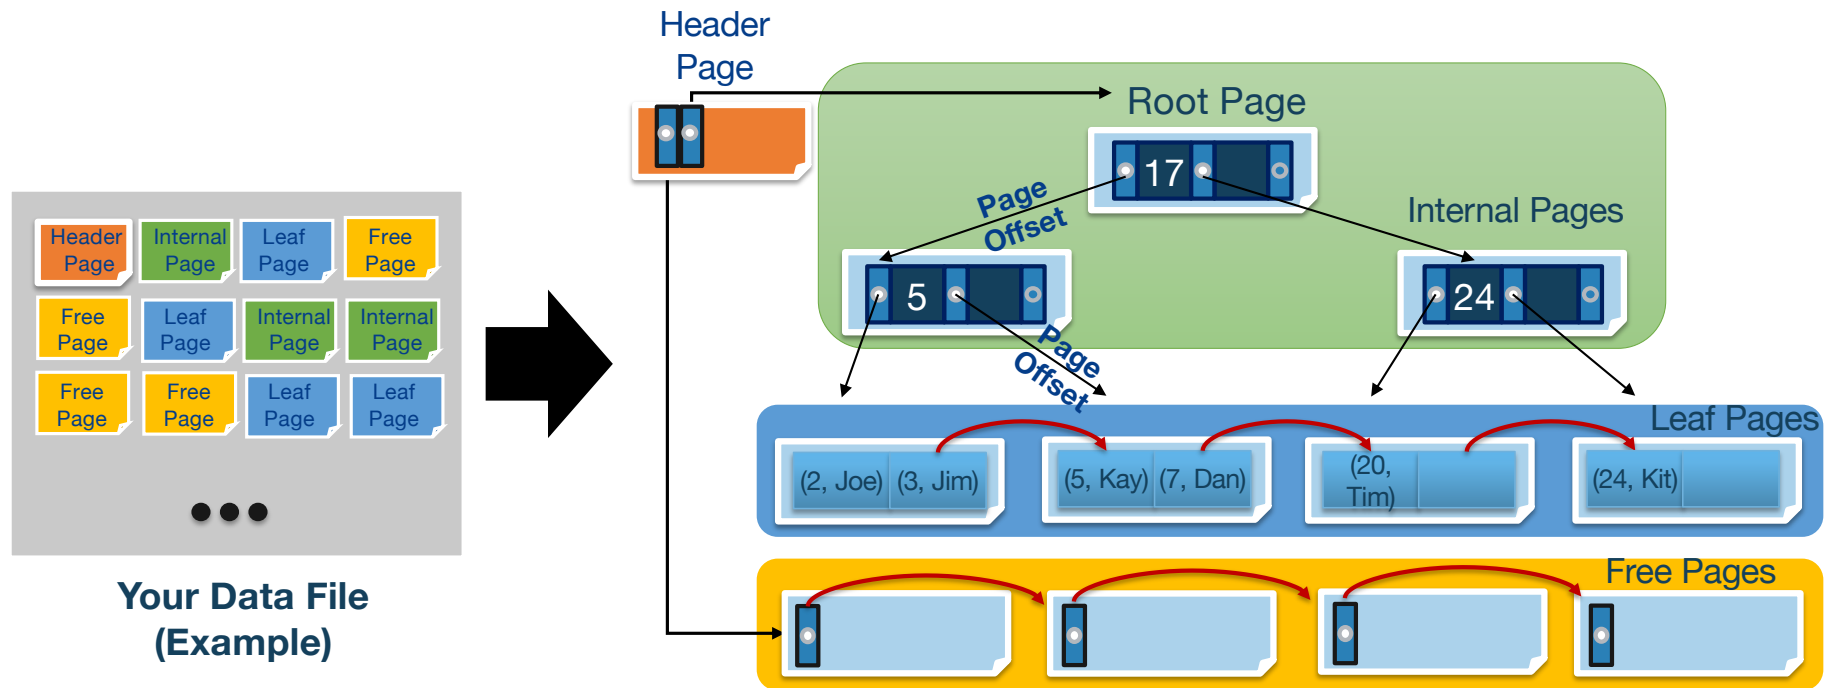
- Implement **close_table** interface.
 - **int close_table(int table_id);**
 - Write all pages of this table from buffer to disk.
 - If success, return 0. Otherwise, return non-zero value.
- Implement database shutdown function.
 - **int shutdown_db();**
 - Flush all data from buffer and destroy allocated buffer.
 - If success, return 0. Otherwise, return non-zero value.

Project Specification

➤ Your library (libbpt.a) should provide those API services.

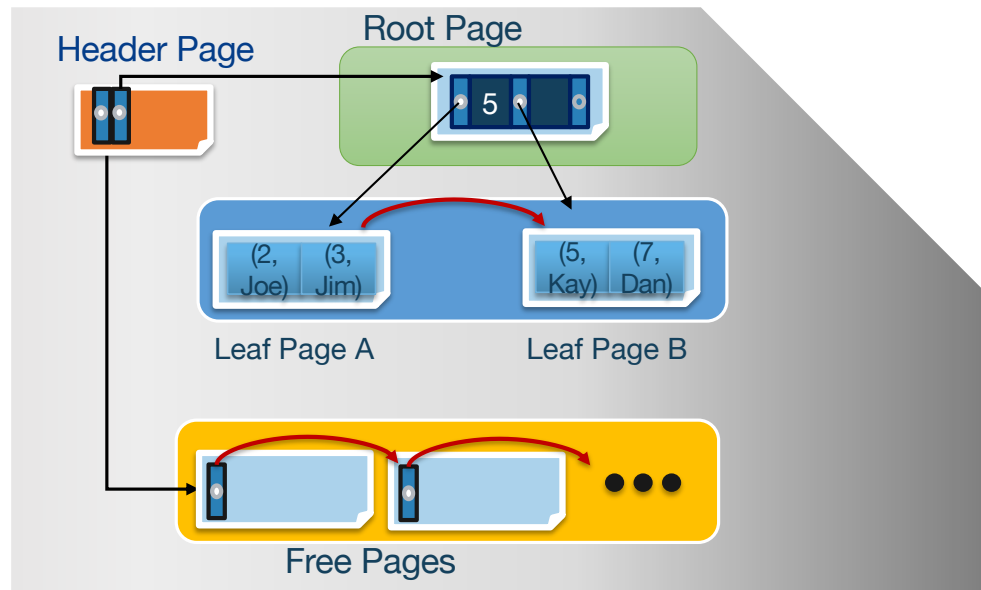
1. **int init_db (int buf_num);**
 - Initialize buffer pool with given number and buffer manager.
2. **int open_table (char * pathname);**
 - Open existing data file using 'pathname' or create one if not existed. If success, return **table_id**.
3. **int db_insert (int table_id, int64_t key, char * value);**
4. **int db_find (int table_id, int64_t key, char* ret_val);**
5. **int db_delete (int table_id, int64_t key);**
6. **int close_table(int table_id);**
 - Write the pages relating to this table to disk and close the table.
7. **int shutdown_db(void);**
 - Destroy buffer manager.

So far..




Buffer Management

- Assume the on-disk pages are stored like below form.



Client

Buffer Manager

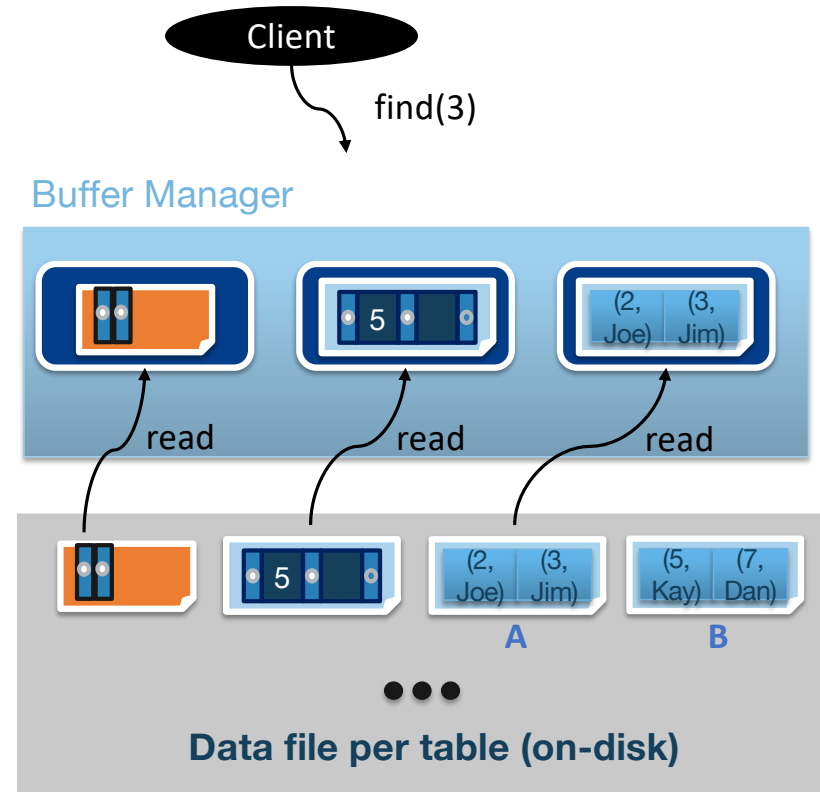


Files and Disk Space Manager



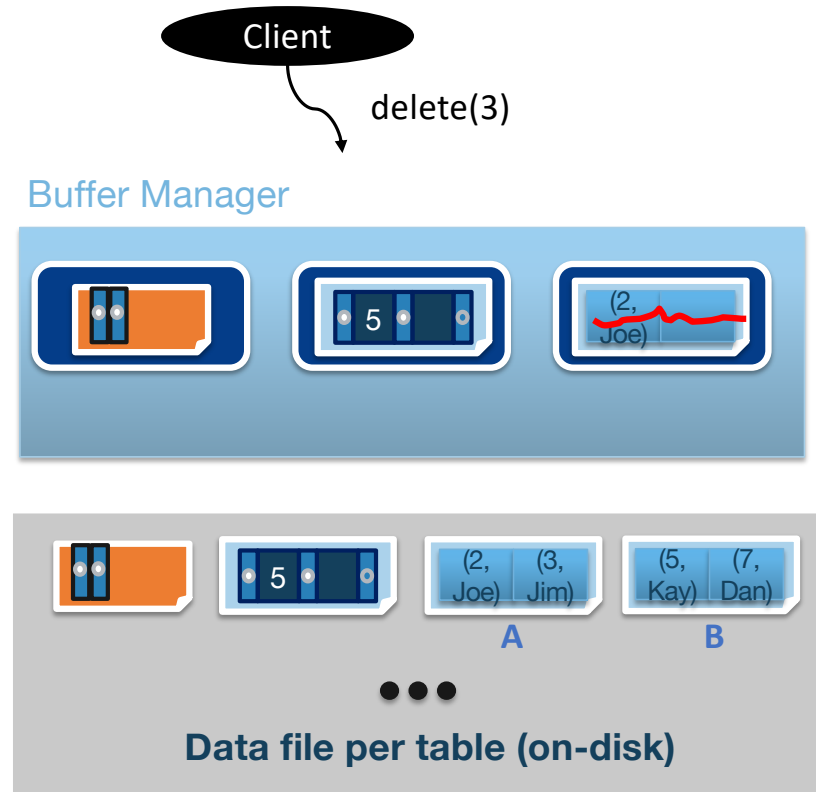
Buffer Management

- First, search the page from the buffer pool.
- If the page is not in the buffer pool (i.e, cache-miss occurs), read the page from disk and maintain this page in a buffer block.
- While indexing from the root to the leaf page A (where key 3 is located), the header page and the root page (internal page) are also read by the buffer manager.



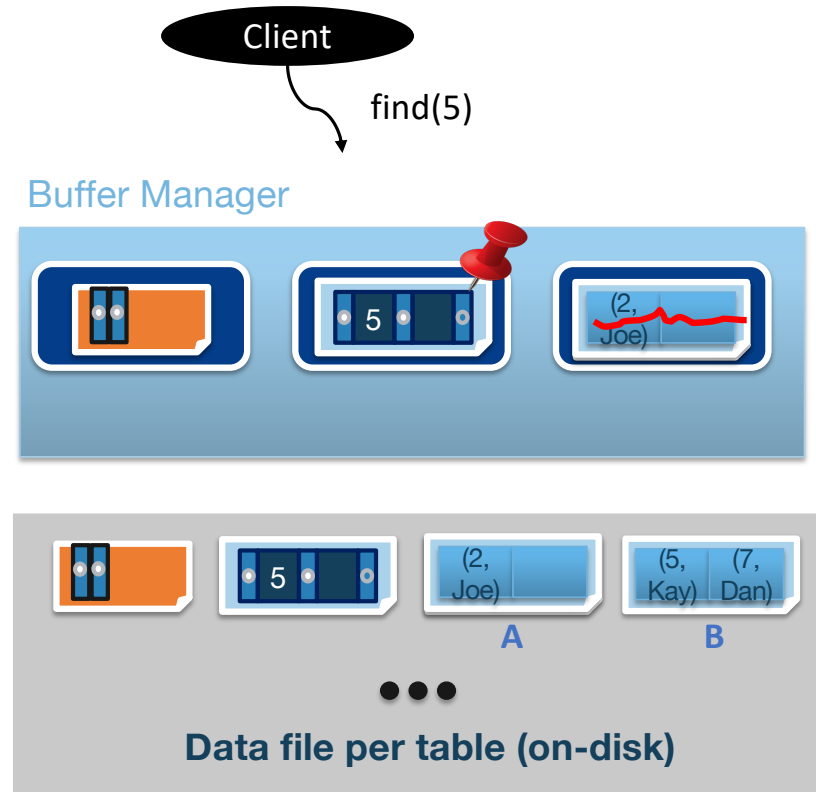
Buffer Management

- After reading the page to the buffer, update operation can be handled in the buffer (memory).
- So “delete key 3” operation occurs in the buffer, which makes that page marked to dirty.



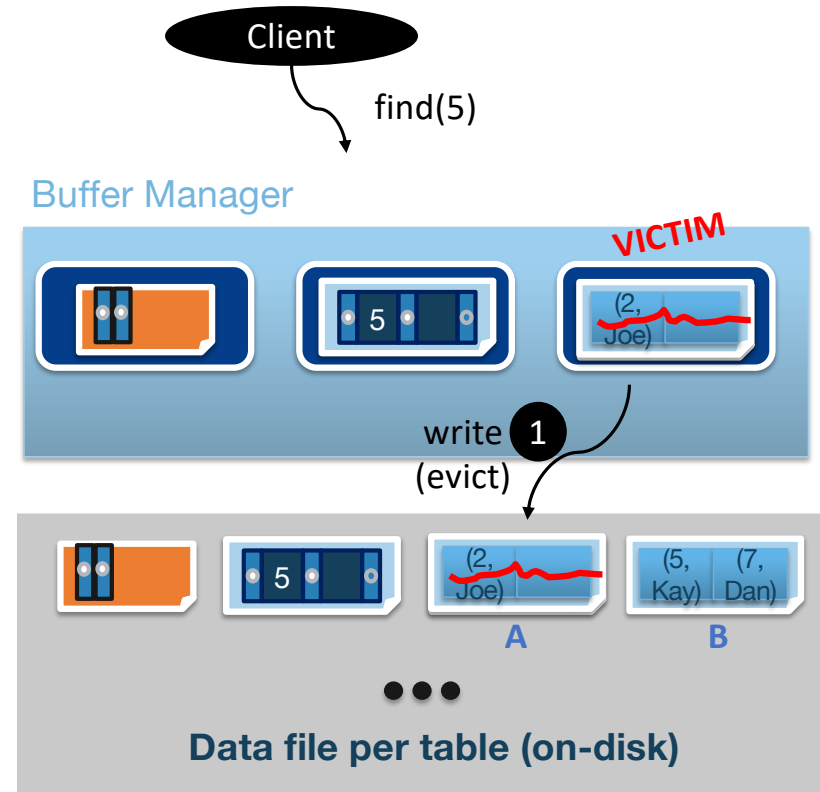
Buffer Management

- Dirty page is written to disk when those page is selected to the victim of LRU policy.
- Assuming the example shown right, find(5) tries to read root page.




Buffer Management

- Dirty page is written to disk when the page is selected to the victim of LRU policy.
- Assuming the example shown right, find(5) tries to read the leaf page B which triggers page eviction. (pinned page should not be the victim of eviction.)
- If the victim page is marked as dirty, write data to disk first. ①



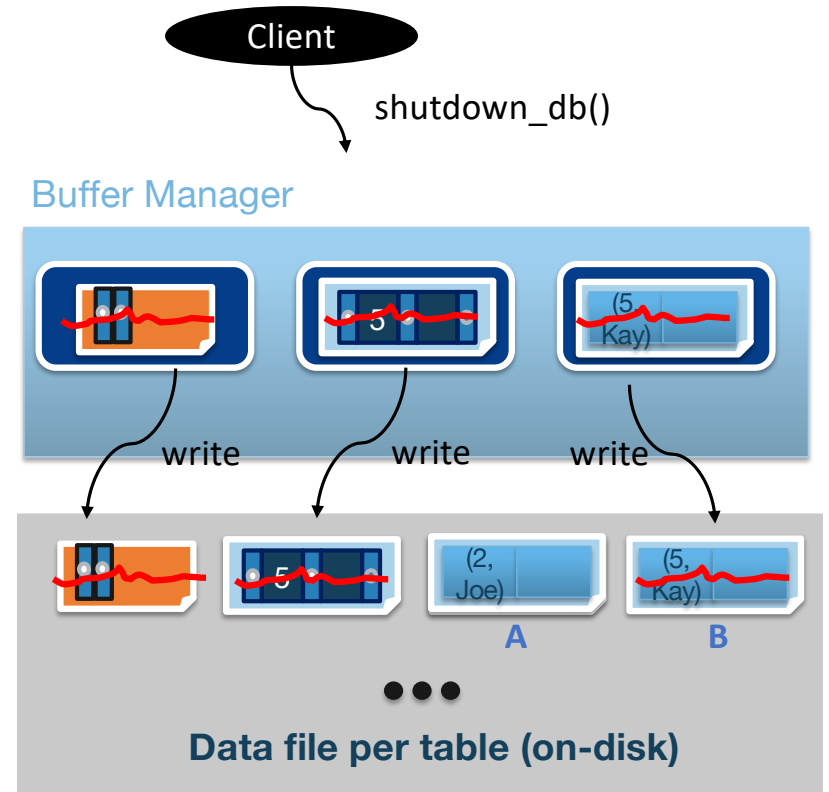
Buffer Management

- Dirty page is written to disk when those page is selected to the victim of LRU policy.
- Assuming the example shown right, find(5) tries to read the leaf page B which triggers page eviction. (pinned page should not be the victim of eviction.)
- If the victim page is marked as dirty, write data to disk first. ❶
- Then read another page from disk. ❷



Buffer Management

- `close_table()` or `shutdown_db()` writes out all dirty buffer block to disk.
- `close_table()` writes out the pages only from those relating to given `table_id`.
- This command can provide synchronous semantic (durability) to the user but lose performance.



File Manager APIs

- The File Manager APIs should be properly changed for Project3.

Submission

- Implement in-memory buffer manager and submit a report (Wiki) including your design
 - Deadline: Nov 1 11:59pm
- We'll only score your commit before the deadline and your submission after the deadline will not be accepted

Thank you
